Tutorial manual: constraint-based modeling for systems medicine

Thierry D.G.A. Mondeel   Stefania Astrologo

Samrina Rehman   Ablikim Abulikemu

Hans V. Westerhoff

May 30, 2016

Contents

The platform: SageMath

We will use the online SageMath platform for this tutorial.

> SageMath is a free open-source mathematics software system licensed under the GPL. It builds on top of many existing open-source packages: NumPy, SciPy, matplotlib, Sympy, Maxima, GAP, FLINT, R and many more. Access their combined power through a

common, Python-based language or directly via interfaces or wrappers. Mission: Creating a viable free open source alternative to Magma, Maple, Mathematica and Matlab.

<div>

ASSIGNMENT

It is best to follow this tutorial in groups of two or three students. This will reduce the server load and allow you to discuss together. Please form a group now if you have not yet.

</div>

<div>

ASSIGNMENT

Go to `https://cloud.sagemath.com/` and register for an account. Then tell the tutorial supervisors the user name of your account so that we can add you to the tutorial project on SageMath.

</div>

A note about doing this on your own laptop

The software required to do constraint-based modeling is not quickly installed. For this reason we utilize the online SageMath platform which already has the required software installed. If you would like to do similar work offline on your own laptop you will require the following things:

- A python distribution with Jupyter notebook installed: we recommend the anaconda scientific python distribution. See: `https://www.continuum.io/downloads`.

- Cameo and Cobrapy will require you to have glpk. See: `https://www.gnu.org/software/glpk/`.

- When you have Anaconda installed you can install Cobrapy and Cameo on top of it. By running the following command (cameo will automatically also install escher and cobrapy)

  ```
  pip install cameo
  ```

  For more details see `https://github.com/biosustain/cameo`.

If you're interested you may try this after today's tutorial.

Another very good online alternative developed at the VU by Brett Olivier is `F-A-M-E.org` which gives you an environment to perform constraint-based modeling without having to code anything.

Setting up your own workspace

Once the tutorial supervisor has added you to the project: the "systems medicine tutorial" project will appear in your list of projects on SageMath.

WATCH OUT: The files in this project are editable by you. We do not want you to edit the original files so you will have to create your own *copy* of them. The sequence of steps to do this is outlined below.

ASSIGNMENT

- Navigate to the main folder of this project. If you are anywhere in the Sagemath website you can get there through the "Projects" button and then selecting the project for this tutorial.

- From the main folder of the project navigate to the "Students" subfolder.

- Click "Create" at the top of the screen.

- Type in your name (this will be the name of your own personal folder) and then click "Folder"

- You will now automatically end up in your own folder, however it will be empty. Lets fix that!

- Navigate back to the main folder by clicking "Parent Directory" twice.

- Check the box next to the "tutorial_files" folder.

- Then click copy and select your personal folder in the drop down menu "Destination" in the top-right of the screen.

- Now navigate to your personal folder again.

You should now have a personal folder with the tutorial files in them. You can safely edit these files to your heart's content!

If anything goes really wrong, you can always copy the original tutorial files back to your personal folder again.

Overview and introduction

Now that you have your own space to work we can start the tutorial. The aim of this tutorial is to introduce you to constraint-based modeling and the applications of the human metabolic reconstruction RECON2 [1].

Note of warning: this is the first time we give this tutorial on this platform so please forgive us for any mistakes/shortcomings. The website might be a bit slow at times, be patient please.

There are various software packages that enable the user to perform constraint-based modeling. We will use the Python based: Cameo `https://github.com/biosustain/cameo` and Cobrapy [2]. This requires you to know at least some basics of python which is what we will start the tutorial with. If you already have experience with Python feel free to skim (not skip) the first part. Even if you already have some experience there may be some tips and tricks in the first part that will come in handy later.

After the Python introduction we will introduce the Cameo package and how to do computational analysis on the human metabolic reconstruction with it. Specifically, we will investigate tyrosine and phenylalanine metabolism and biomarker predictions for inborn errors of metabolism like phenylketonuria.

Happy computing!

1 Introduction to Python & Jupyter

To do our computational analysis on the human metabolic reconstruction we will make use of the Python programming language.

Some exerpts from the wiki page on Python `https://en.wikipedia.org/wiki/Python_(programming_language)`

> Python is a widely used high-level, general-purpose, interpreted, dynamic programming language. Its design philosophy emphasizes code readability, and its syntax allows programmers to express concepts in fewer lines of code than would be possible in languages such as C++ or Java. The language provides constructs intended to enable clear programs on both a small and large scale.

> Python was conceived in the late 1980s, and its implementation began in December 1989 by Guido van Rossum at Centrum Wiskunde & Informatica (CWI) in the Netherlands as a successor to the ABC language (itself inspired by SETL) capable of exception handling and interfacing with the operating system Amoeba. Van Rossum is Python's principal author, and his continuing central role in deciding the direction of Python is reflected in the title given to him by the Python community, benevolent dictator for life (BDFL).

> The core philosophy of the language is summarized by the document The Zen of Python (PEP 20), which includes aphorisms such as:
> - Beautiful is better than ugly
> - Explicit is better than implicit
> - Simple is better than complex
> - Complex is better than complicated
> - Readability counts

We will write simple python code to do constraint-based modeling on the human metabolic map using so-called Jupyter notebooks:

Notebook documents (or notebooks, all lower case) are documents produced by the Jupyter Notebook App which contain both computer code (e.g. python) and rich text elements (paragraph, equations, figures, links, etc...). Notebook documents are both human-readable documents containing the analysis description and the results (figures, tables, etc..) as well as executable documents which can be run to perform data analysis.

We have generated several such notebooks for you. These contain ever more complex commands and analysis and build on one another. Each notebook contains some examples of how to do specific things and one or more small questions and assignments.

Getting to know the Jupyter notebook

ASSIGNMENT

All Jupyter notebooks for this tutorial are located in the "tutorial_files" folder within your personal folder. We will go through them in their numbered order, so your first Jupyter notebook is "0_running_code_in_notebook.ipynb" etc.
Open the first file and go through the various commands within in. It will introduce you to the basics of the notebook. Make sure you interact with the document. Execute code cells manually, don't just look at it, and do the assignments.
This part should not take more than 15 minutes.

Python essentials

Now that we know our way around the Jupyter notebook it's time to become a bit more familiar with Python itself so that you are able to to do computational analysis later on.

ASSIGNMENT

Go through tutorial 1: "1_introduction_to_python.ipynb" until the section "Functions". The material from functions onward is not needed to understand the rest of the tutorial. If you have time to spare later on, you can go back to this tutorial.
Again, if you are familiar with Python already, feel free to skim this tutorial.
This tutorial should take you about 30 min. to go through.

2 Introduction to CAMEO

Now that we understand some basic Python syntax we can get into the specific python packages that allow us to do constraint-based modeling analyses on the human metabolic map. The following tutorial will introduce the Cameo package.

> `http://cameo.bio/` (Developed at the Novo Nordisk Foundation Center for Biosustainability)
>
> Cameo is a high-level python library developed to aid the strain design process in metabolic engineering projects. The library provides a modular framework of simulation methods, strain design methods, access to models, that targets developers that want custom analysis workflows.

The cameo website showcases some basic Jupyter notebooks on how the use the package.

> **ASSIGNMENT**
>
> Visit `http://try.cameo.bio/`. This opens up a special notebook server just for you. You should see a list of tutorials. We will be doing tutorial 1 (quick start), 3 (simulate models) and 4 (analyze models).
> These notebooks do not contain any questions so it is imperative that you make sure yourself that you understand the principles showcased in the notebooks. If you run into any unknown terms ask the tutorial supervisors or ask Google. In the assignment boxes below we added some questions to keep in mind while going through these tutorials.

Crash course on Cameo

> **ASSIGNMENT**
>
> Run through tutorial 1 (Quick start). This tutorial will introduce you to the bare essentials of the Cameo package.
> **Some self-check questions:**
>
> - use the dir() function to figure out attributes that various model objects have. Try dir() on the model variable, the metabolic glucose, and one of the reactions. Especially focus on the ones without underscores.

Flux balance analysis

We will use the constraint-based modeling technique called flux balance analysis (FBA) [3]. Given the metabolic network this method computationally calculates the flux distribution, meaning the fluxes through all reactions in the model at

steady-state, that is optimal according to a specified objective function. In our case, the objective function will usually be the flux through the biomass reaction. This reaction contains (when known) experimentally determined ratio's of metabolites that make of the composition of a cell.

FBA has traditionally been applied to metabolism as we are doing here. FBA primarily makes use of the stoichiometry matrix S. This matrix is of size $m \times r$ where m is the number of species or metabolites in the system and r is the number of reactions. Every row of S specifies for a specific metabolite in what quantity it partakes in each reaction. Therefore, each element (i, j) of S contains the stoichiometric coefficient of metabolite i in reaction j.

The word flux is used to describe a reaction rate at steady state. In FBA the aim is to find the flux distributions of the network that satisfy (1) the steady-state condition, this is based on the assumption that metabolism occurs on a fast time-scale compared to gene regulatory events and thus reaction rates are constant [4], (2) thermodynamic feasibility , i.e. some reactions are known to be irreversible, (3) maximal flux constraints when these are known and (4) maximize a linear function of the fluxes. Mathematically, this can be concisely summarized as

$$\text{Maximize } Z = c^T v, \text{ such that}$$
$$Sv = 0$$
$$\alpha_k \leq v_k \leq b_k, \text{ for all k.}$$

Here v is the vector of fluxes representing all reactions in the model, α_k is the lower bound for reaction k and similarly β_k is the upper bound for reaction k.

This can be thought of as first constraining all possible solutions to the ones that allow a steady state and satisfy the bounds (this results in a multi dimensional cone within the null space of S) and then finding the optimal solution among the remaining degrees of freedom.

We will perform in-silico experiments with genetic manipulations by changing the third line in the above equation, i.e. the lower and upper bounds for the reactions. By setting both the lower bound α_k and the upper bound β_k to zero we in effect knock out a reaction. Similarly, by setting it to a specific value we can fix a certain flux level for a specific reaction.

Flux variability analysis

In contrast, in flux variability analysis [5, 6] a linear programming problem is run for each reaction in the map. First, FVA runs FBA and determines the maximal value of the objective function. Then, for any reaction j in the network, the objective is set first to minimizing the flux and then to maximizing the flux while satisfying the same requirements as FBA: steady- state, thermodynamics and upper/lower bounds and additionally forcing a specified percentage of the optimal flux through the objective flux(es). FVA can mathematically be summarized as

$$\text{Maximize } v_i, \text{ such that}$$
$$Sv = 0$$
$$Z \geq pZ_{opt}$$
$$\alpha_k \leq v_k \leq b_k, \text{ for all k.}$$

where Z_{opt} is the optimal value of the objective function identified by FBA and $0 \leq p \leq 1$ is the minimal percentage of this optimum that should be achieved with FVA. For the reaction j this then gives a range of fluxes that are consistent with the range of values allowed for the objective function, when allowing all other fluxes to vary arbitrarily. This process can be repeated for all reactions j, leading to a different range of flux distributions for each reaction. If the network has r reactions, one thereby obtains r ranges for flux distributions. Applications of FVA range from exploration of optimal flux distributions other than the one identified by FBA, through studying flux distributions under suboptimal performance, to investigating network flexibility or redundancy.

3   Introduction to RECON2

Using our newly acquired cameo knowledge we now turn to the human metabolic
reconstruction. For reference see: `http://www.nature.com/nbt/journal/v31/n5/full/nbt.2488.html`.

We now go back to the SageMath project.

4 Phenylketonuria: an introduction

In the last tutorial we began to look at the enzyme catalyzing the reaction from phenylalanine to tyrosine. We will study it further now, specifically when it is defective and what the consequences of that are.

Phenylalanine Catabolism Is Genetically Defective in Some People

The following is taking from Lehninger's principles of biochemistry [7]:

> Given that many amino acids are either neuro-transmitters or precursors or antagonists of neutrotransmitters, genetic defects of amino acid metabolism can cause defective neural development and mental retardation. In most such diseases specific intermediates accumulate. For example, a genetic defect in phenylalanine hydroxylase, the first enzyme in the catabolic pathway for phenylalanine, is responsible for the disease phenylketonuria (PKU), the most common cause of elevated levels of phenylalanine (hyperphenylalaninemia). Phenylalanine hydroxylase (also called phenylalanine-4-monooxygenase) is one of a general class of enzymes called mixed-function oxidases, all of which catalyze simultaneous hydroxylation of a substrate by an oxygen atom of O_2 and reduction of the other oxygen atom to H_2O. Phenylalanine hydroxylase requires the cofactor tetrahydrobiopterin, which carries electrons from NADH to O_2 and becomes oxidized to dihydrobiopterin in the process. It is subsequently reduced by the enzyme dihydrobiopterin reductase in a reaction that requires NADH.

> In individuals with PKU, a secondary, normally little-used pathway of phenylalanine metabolism comes into play. In this pathway phenylalanine undergoes transamination with pyruvate to yield

phenylpyruvate. Phenylalanine and phenylpyruvate accumulate in the blood and tissues and are excreted in the urinehence the name phenylketonuria. Much of the phenylpyruvate, rather than being excreted as such, is either decarboxylated to phenylacetate or reduced to phenyllactate. Phenylacetate imparts a characteristic odor to the urine, which nurses have traditionally used to detect PKU in infants. The accumulation of phenylalanine or its metabolites in early life impairs normal development of the brain, causing severe mental retardation. This may be caused by excess phenylalanine competing with other amino acids for transport across the blood-brain barrier, resulting in a deficit of required metabolites.

Phenylketonuria was among the first inheritable metabolic defects discovered in humans. When this condition is recognized early in infancy, mental retardation can largely be prevented by rigid dietary control. The diet must supply only enough phenylalanine and tyrosine to meet the needs for protein synthesis. Consumption of protein-rich foods must be curtailed. Natural proteins, such as casein of milk, must first be hydrolyzed and much of the phenylalanine removed to provide an appropriate diet, at least through childhood. Because the artificial sweetener aspartame is a dipeptide of aspartate and the methyl ester of phenylalanine, foods sweetened with aspartame bear warnings ad- dressed to individuals on phenylalanine-controlled diets.

Also see Figure 1 from [7] below for a pathway visualization of some more common IEMs.

Phenylketonuria can also be caused by a defect in the enzyme that catalyzes the regeneration of tetrahydrobiopterin. The treatment in this case is more complex than restricting the intake of phenylalanine and tyrosine.

Tetrahydrobiopterin is also required for the formation of L-3,4 dihydroxyphenylalanine (L-dopa) and 5-hydroxytryptophan precursors of the neurotransmitters norepinephrine and serotonin, respectively and in phenylketonuria of this type, these pre- cursors must be supplied in the diet. Supplementing the diet with tetrahydrobiopterin itself is ineffective because it is unstable and does not cross the blood-brain barrier.

Also see Figure 2 from [7] for the salvage reaction of biopterin.

Figure 1: Catabolic pathways for phenylalanine and tyrosine. In humans these amino acids are normally converted to acetoacetyl-CoA and fumarate. Genetic defects in many of these enzymes cause inheritable human diseases (shaded yellow).


FIGURE 18–24 **Role of tetrahydrobiopterin in the phenylalanine hydroxylase reaction.** The H atom shaded pink is transferred directly from C-4 to C-3 in the reaction. This feature, discovered at the NIH, is called the NIH Shift.

Figure 2: Role of tetrahydrobiopterin in the phenylalanine hydroxylase reaction.

5 Biomarker prediction using flux variability analysis: intro

Next we will focus on an intriguing application of the human metabolic model:
the ability to predict the accumulation and depletion of any biomarker metabolites in extracellular tissues and fluids upon a specific mutation.

The method we will utilize to do this is based on flux variability analysis
and was presented originally in [8]. It was subsequently used in the RECON2
paper to produce Figure 3.

6 Phenylketonuria: biomarker prediction

Now that you have a feel for the method developed by Shlomi et al, we are
going to apply it specifically to the PKU case. Note that this IEM was absent
from Figure 3 in the RECON2 paper. It was however present in the article by
Shlomi et al, but there the predecessor of RECON2, RECON1, was used for the
prediction.

7 Concluding remarks

We hope this tutorial achieved its aims to:

- Introduce you to programming with Python

- Show you some real life examples of using metabolic maps and constraint-based modeling

- Illustrate how genome-wide metabolic maps may be used to go beyond "textbook understanding" of pathways and show novel unexpected alternative ways to achieve metabolic tasks.

- Highlighted one interesting way to use metabolic maps for biomarker prediction in tissues and biofluids.

For more information see the references below.

References

[1] I. Thiele, N. Swainston, R. M. T. Fleming, A. Hoppe, S. Sahoo, M. K. Aurich, H. Haraldsdottir, M. L. Mo, O. Rolfsson, M. D. Stobbe, S. G. Thorleifsson, R. Agren, C. Bölling, S. Bordel, A. K. Chavali, P. Dobson, W. B. Dunn, L. Endler, D. Hala, M. Hucka, D. Hull, D. Jameson, N. Jamshidi, J. J. Jonsson, N. Juty, S. Keating, I. Nookaew, N. Le Novère, N. Malys, A. Mazein, J. A. Papin, N. D. Price, E. Selkov, M. I. Sigurdsson, E. Simeonidis, N. Sonnenschein, K. Smallbone, A. Sorokin, J. H. G. M. van Beek, D. Weichart, I. Goryanin, J. Nielsen, H. V. Westerhoff, D. B. Kell, P. Mendes, and B. Ø. Palsson, "A community-driven global reconstruction of human metabolism.," *Nature biotechnology*, vol. 31, no. 5, pp. 419–425, 2013.

[2] A. Ebrahim, J. A. Lerman, B. O. Palsson, and D. R. Hyduke, "COBRApy: COnstraints-Based Reconstruction and Analysis for Python," *BMC Systems Biology*, vol. 7, no. 1, p. 74, 2013.

[3] J. D. Orth, I. Thiele, and B. Ø. Palsson, "What is flux balance analysis?," *Nature biotechnology*, vol. 28, no. 3, pp. 245–248, 2010.

[4] R. Heinrich and S. Schuster, "The regulation of cellular systems," *Book*, p. 416, 1996.

[5] S. Gudmundsson and I. Thiele, "Computationally efficient flux variability analysis.," *BMC bioinformatics*, vol. 11, no. 1, p. 489, 2010.

[6] R. Mahadevan and C. H. Schilling, "The effects of alternate optimal solutions in constraint-based genome-scale metabolic models," *Metabolic Engineering*, vol. 5, no. 4, pp. 264–276, 2003.

[7] A. Lehninger, D. L. Nelson, and M. M. Cox, *Lehninger Principles of Biochemistry*. W. H. Freeman, fifth edition ed., June 2008.

[8] T. Shlomi, M. N. Cabili, and E. Ruppin, "Predicting metabolic biomarkers of human inborn errors of metabolism," *Molecular Systems Biology*, vol. 5, no. 263, p. 263, 2009.